

Automated Verification of Systems Code using Type-Based Memory Abstractions

Olivier Nicole

19 April, 2022

École normale supérieure – PSL, CEA List



Alan Schmitt
Mihaela Sighireanu
Timothy Bourke
Jean-Christophe Filliâtre
Julia Lawall
Xavier Rival
Matthieu Lemerre

Inria
ENS Paris-Saclay
Inria
CNRS
Inria
ENS, Inria
CEA List

Reviewer
Reviewer
Examiner
Examiner
Examiner
Supervisor
Co-advisor


Software correctness in critical systems




Software contains bugs and vulnerabilities

New bugs are discovered regularly

CVE-2021-22555 in **Linux**




Kernels are a critical component of software systems




- The operating system kernel:
 - organizes the sharing of hardware resources between applications
 - prevents applications from disturbing the functioning of other applications

Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...
- **Privilege escalation**


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...
- **Privilege escalation**


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...
- **Privilege escalation**


Kernels are a critical component of software systems



Two bugs particularly threaten safety and security:

- **Runtime errors** Division by zero, illegal memory access...
- **Privilege escalation**
- **Goal:** Verify their absence using formal methods

Kernel execution: the system loop




A system with a kernel executes alternatively **user code** and **kernel runtime**.

It may **pause** a task and run another one

```
struct Task {  
    Memory_table *mem_table;  
    Context ctx;  
    Task *next;  
};
```

Memory protection

- The **Memory Protection Unit (MPU)** controls access to memory
- Restricted to an **interval of memory addresses**
- Interval described in a **memory protection table**



- mpu can be changed only if `privilege = 1`

Kernel verification requires memory invariants

```
Task *cur;
```

```
void handle_timer() {  
    /* Save task context */  
    cur→ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur→next;  
    /* Load new memory protection */  
    mpu = cur→mem_table;  
    /* Give control to new task */  
    switch_context(&cur→ctx);  
}
```

```
struct Task {  
    Memory_table *mem_table;  
    Context ctx;  
    Task *next;  
};
```

Kernel verification requires memory invariants

```
Task *cur;
```

```
void handle_timer() {  
    /* Save task context */  
    cur→ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur→next;  
    /* Load new memory protection */  
    mpu = cur→mem_table;  
    /* Give control to new task */  
    switch_context(&cur→ctx);  
}
```

```
struct Task {  
    Memory_table *mem_table;  
    Context ctx;  
    Task *next;  
};
```

Invariants needed

- cur does not point to a memory table

Kernel verification requires memory invariants

```
Task *cur;
```

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

```
struct Task {  
    Memory_table *mem_table;  
    Context ctx;  
    Task *next;  
};
```

Invariants needed

- `cur` does not point to a memory table
- at the end,
 `mpu->start > kernel_last_addr`

Kernel verification requires memory invariants

```
Task *cur;
```

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

```
struct Task {  
    Memory_table *mem_table;  
    Context ctx;  
    Task *next;  
};
```

Invariants needed

- `cur` does not point to a memory table
- at the end,
`mpu->start > kernel_last_addr`
`cur->ctx.privilege = 0`

Invariant proof methods for kernel verification

Interactive proof

- seL4 [SOSP'09]
- CertiKOS [OSDI'16]

Deductive verification

- Verve [PLDI'10]
- Komodo [SOSP'17]

Prove strong properties, but require a lot of work from experts

Invariant proof methods for kernel verification

Interactive proof

- seL4 [SOSP'09]
- CertiKOS [OSDI'16]

Deductive verification

- Verve [PLDI'10]
- Komodo [SOSP'17]

Prove strong properties, but require a lot of work from experts

“Push-button” verification

- PROSPER [CCS'13]
- Serval [SOSP'19]
- Phidias [EuroSys'20]

- Still require to write kernel invariants
- Only support *bounded loops* (no priority scheduling)
- Requires a *fixed memory layout* (depends on the number of tasks)

Invariant proof methods for kernel verification

Interactive proof

- seL4 [SOSP'09]
- CertiKOS [OSDI'16]

Deductive verification

- Verve [PLDI'10]
- Komodo [SOSP'17]

Prove strong properties, but require a lot of work from experts

“Push-button” verification

- PROSPER [CCS'13]
- Serval [SOSP'19]
- Phidias [EuroSys'20]

- Still require to write kernel invariants
- Only support *bounded loops* (no priority scheduling)
- Requires a *fixed memory layout* (depends on the number of tasks)

Sound static analysis by abstract interpretation

- ASTERIOS

- Applied to the **kernel executable**
- Infers invariants
- Handles unbounded loops
- Handles parameterized verification
- Low annotation burden (e.g. 58 lines)

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x,y,z,r;  
x = random(1, 4);  
y = random(1, 4);  
  
z = x - y;  
r = z * z;  
  
assert (r <= 9);  
assert (r >= 0);
```

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x,y,z,r;
```

```
x = random(1, 4);
```



$x \in [1, 4]$

```
y = random(1, 4);
```

```
z = x - y;
```

```
r = z * z;
```

```
assert (r <= 9);
```

```
assert (r >= 0);
```

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x, y, z, r;
```

```
x = random(1, 4);    ○—————  $x \in [1, 4]$ 
```

```
y = random(1, 4);    ○—————  $y \in [1, 4]$ 
```

```
z = x - y;
```

```
r = z * z;
```

```
assert (r <= 9);
```

```
assert (r >= 0);
```

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x, y, z, r;
```

```
x = random(1, 4);    ○—————  $x \in [1, 4]$ 
```

```
y = random(1, 4);    ○—————  $y \in [1, 4]$ 
```

```
z = x - y;    ○—————  $z \in [-3, 3]$ 
```

```
r = z * z;
```

```
assert (r <= 9);
```

```
assert (r >= 0);
```


Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x, y, z, r;
```

```
x = random(1, 4);    ○—————  $x \in [1, 4]$ 
```

```
y = random(1, 4);    ○—————  $y \in [1, 4]$ 
```

```
z = x - y;    ○—————  $z \in [-3, 3]$ 
```

```
r = z * z;    ○—————  $r \in [-9, 9]$ 
```

```
assert (r <= 9);
```

```
assert (r >= 0);
```

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x, y, z, r;
```

```
x = random(1, 4);   ○—————  $x \in [1, 4]$ 
```

```
y = random(1, 4);   ○—————  $y \in [1, 4]$ 
```

```
z = x - y;   ○—————  $z \in [-3, 3]$ 
```

```
r = z * z;   ○—————  $r \in [-9, 9]$ 
```

```
assert (r <= 9);   ○————— True
```

```
assert (r >= 0);
```

Sound static analysis by abstract interpretation

Illustration: Verifying numerical assertions

```
int x, y, z, r;
```

```
x = random(1, 4);    ○—————  $x \in [1, 4]$ 
```

```
y = random(1, 4);    ○—————  $y \in [1, 4]$ 
```

```
z = x - y;    ○—————  $z \in [-3, 3]$ 
```

```
r = z * z;    ○—————  $r \in [-9, 9]$ 
```

```
assert (r <= 9);    ○————— True
```

```
assert (r >= 0);    ○————— Maybe
```

Invariants inferred by pointer analyses

- Infer **points-to** or **aliasing** relations between program variables

$x \mapsto \{\&y, \&z\}$ “x can pointer to either y or z”

- Very **efficient**
- Difficult to apply to machine code
- Do not consider the **structure** of data nor **values** in memory

Invariants inferred by pointer analyses

Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur→ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur→next;  
    /* Load new memory protection */  
    mpu = cur→mem_table;  
    /* Give control to new task */  
    switch_context(&cur→ctx);  
}
```

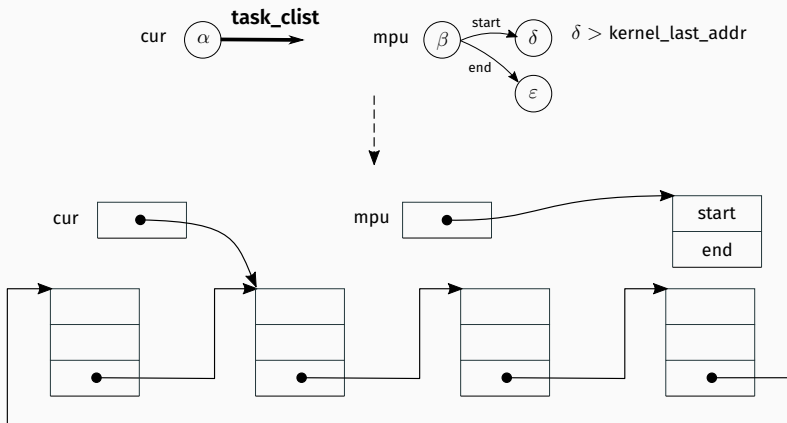
$cur \mapsto ?$

$mpu \mapsto ?$

Shape analyses

Shape analyses are good candidates...

- Verify **strong properties** on memory



... but difficult to apply

- Difficult to apply to low-level code
- Requires **case disjunctions** which can be costly
- Fragile in presence of **precision loss**

Type-based memory invariants

Physical types describe the layout of values in memory

```
type Memory_table = {x : int | x > kernel_last_addr} × int
```

| |
|-------|
| start |
| end |



- **Product**: contiguous values in memory
- **Refinement types**: numerical constraints on values

Type-based memory invariants

Physical types describe the layout of values in memory

type **Memory_table** = $\{x : \text{int} \mid x > \text{kernel_last_addr}\} \times \text{int}$

type **Context** = $\dots \times \{\text{privilege} : \text{int} \mid \text{privilege} = 0\} \times \dots$



- **Product**: contiguous values in memory
- **Refinement types**: numerical constraints on values


Type-based memory invariants

Physical types describe the layout of values in memory

type **Memory_table** = $\{x : \text{int} \mid x > \text{kernel_last_addr}\} \times \text{int}$


type **Context** = $\dots \times \{\text{privilege} : \text{int} \mid \text{privilege} = 0\} \times \dots$

type **Task** = **Memory_table**.(0)* \times **Context** \times $\{x : \text{Task}.\text{(0)}^* \mid x \neq 0\}$





- **Product**: contiguous values in memory
- **Refinement types**: numerical constraints on values
- **Pointer types** with offset


Separation



- Tasks not necessarily separated



Separation



- **Tasks** not necessarily separated
- **Task** and **Context** may alias because **Task** contains a **Context**
- But **Task** and **Memory_table** never alias:

If $\alpha : \text{Task}.(0)^*$ and $\beta : \text{Memory_table}.(0)^*$,
then $\alpha \neq \beta$.

Type-based memory invariants

```
type Memory_table = {x: int | x > kernel_last_addr} × int  
type Context = ... × {privilege: int | privilege = 0} × ...  
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

```
void handle_timer() {  
    /* Save task context */  
    cur→ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur→next;  
    /* Load new memory protection */  
    mpu = cur→mem_table;  
    /* Give control to new task */  
    switch_context(&cur→ctx);  
}
```

Type-based memory invariants

```
type Memory_table = {x: int | x > kernel_last_addr} × int  
type Context = ... × {privilege: int | privilege = 0} × ...  
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

$\alpha \neq 0$

```
void handle_timer() { ○———— cur  $\alpha : \text{Task}.(0)^*$  mpu  $\beta : \text{Memory\_table}.(0)^*$   
  /* Save task context */  
  cur→ctx = interrupted_ctx;  
  /* Schedule next task */  
  cur = cur→next;  
  /* Load new memory protection */  
  mpu = cur→mem_table;  
  /* Give control to new task */  
  switch_context(&cur→ctx);  
}
```

Type-based memory invariants

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

$\alpha \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

○ ————— cur $\alpha : \text{Task}.(0)^*$ mpu $\beta : \text{Memory_table}.(0)^*$

○ ————— cur $\alpha : \text{Task}.(0)^*$ mpu $\beta : \text{Memory_table}.(0)^*$

Type-based memory invariants

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\delta \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

Diagram illustrating memory invariants for the `handle_timer` function:

- Initial state: `cur` points to a task with invariant $\alpha : \text{Task}.(0)^*$, and `mpu` points to a memory table with invariant $\beta : \text{Memory_table}.(0)^*$.
- After `cur→ctx = interrupted_ctx;`: `cur` still points to the same task with invariant $\alpha : \text{Task}.(0)^*$, and `mpu` still points to the same memory table with invariant $\beta : \text{Memory_table}.(0)^*$.
- After `cur = cur→next;`: `cur` points to a new task with invariant $\delta : \text{Task}.(0)^*$, and `mpu` still points to the same memory table with invariant $\beta : \text{Memory_table}.(0)^*$.

Type-based memory invariants

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\delta \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

| | | | | |
|--------|-----|------------------------------|-----|--|
| ○————— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○————— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |

Type-based memory invariants


```
type Memory_table = {x: int | x > kernel_last_addr} × int  
type Context = ... × {privilege: int | privilege = 0} × ...  
type Task = Memory_table.(0)* × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

| | | | | | |
|---|--------|-----|------------------------------|-----|--|
| <code>void handle_timer() {</code> | ○————— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| <code>/* Save task context */</code> | | | | | |
| <code>cur→ctx = interrupted_ctx;</code> | ○————— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| <code>/* Schedule next task */</code> | | | | | |
| <code>cur = cur→next;</code> | ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| <code>/* Load new memory protection */</code> | | | | | |
| <code>mpu = cur→mem_table;</code> | ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |
| <code>/* Give control to new task */</code> | | | | | |
| <code>switch_context(&cur→ctx);</code> | ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |
| <code>}</code> | | | | | |

- A **type-based memory abstraction**;
 - Structural invariants on memory encoded as type safety
 - Precision improvements using points-to predicates
- Two **low-level type-based program analyses**
 - For C and machine code
- A method to verify **absence of run-time errors (ARTE)** and **absence of privilege escalation (APE)** on embedded kernels
 - Applied to an unmodified industrial kernel


Tagging memory with types




Tagging memory with types

```
typedef int data;  
struct list {  
    list *next;  
    data d;  
};
```

```
type list = list.(0)* × data  
type data = word4
```




Type-based analysis: Overview



$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0xcafe \end{array} \right]$$

Start from the untyped semantics...

Type-based analysis: Overview



$$n = *l \Rightarrow$$




$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0xcafe \end{array} \right]$$


$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0x100c \end{array} \right]$$

Start from the untyped semantics...


Type-based analysis: Overview



$$n = * \mathfrak{l} \Rightarrow$$



$$*n = \mathfrak{l} \Rightarrow$$



$$\left[\begin{array}{l} \mathfrak{l} \mapsto 0x100c \\ n \mapsto 0xcafe \end{array} \right]$$


$$\left[\begin{array}{l} \mathfrak{l} \mapsto 0x100c \\ n \mapsto 0x100c \end{array} \right]$$

$$\left[\begin{array}{l} \mathfrak{l} \mapsto 0x100c \\ n \mapsto 0x100c \end{array} \right]$$

Start from the untyped semantics...


Type-based analysis: Overview

$$\left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{word}_4 \end{array} \right]$$




$$n = *l \Rightarrow$$

$$\left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{list}.(0)* \end{array} \right]$$



$$*n = l \Rightarrow$$

$$\left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{list}.(0)* \end{array} \right]$$



$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0xcafe \end{array} \right]$$

$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0x100c \end{array} \right]$$

$$\left[\begin{array}{l} l \mapsto 0x100c \\ n \mapsto 0x100c \end{array} \right]$$

Add types to the heap and variables...

Type-based analysis: Overview

$$\left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{word}_4 \end{array} \right] \quad n = *l \quad \Rightarrow \quad \left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{list}.(0)* \end{array} \right] \quad *n = l \quad \Rightarrow \quad \left[\begin{array}{l} l \mapsto \text{list}.(0)* \\ n \mapsto \text{list}.(0)* \end{array} \right]$$

$$\gamma(\text{Env}) = \{s \mid s \text{ is well typed with environment Env}\}$$

Keep only variable types; memory is abstracted by type invariants.

Combination with numerical predicates

- Necessary for:
 - Non-nullness of pointers
 - Array indices and pointer arithmetic
 - Use and verify refinement predicates

Reduced product with a numerical abstract domain

$$\begin{array}{ll} \mathfrak{l} & \boxed{\lambda : \text{list}.(0)*} \quad \lambda \neq 0 \\ \mathfrak{n} & \boxed{\eta : \text{word}_4} \quad \eta \leq 10 \end{array}$$

Example analysis

```
type data = word4  
type list =  
  list.(0)* × data
```

```
if (l != 0) {  
  n = *l;  
}
```

l $\lambda : \text{list}.(0)^*$

n $\eta : \text{word}_4$

Example analysis

```
type data = word4
type list =
  list.(0)* × data
```

```
if (l != 0) {
  n = *l;
}
```

l $\lambda : \text{list}.(0)^*$

n $\eta : \text{word}_4$

↓

assume(l != 0)

l $\lambda : \text{list}.(0)^*$

$\lambda \neq 0$

n $\eta : \text{word}_4$

Example analysis

```
type data = word4
type list =
  list.(0)* × data
```

```
if (l != 0) {
  n = *l;
}
```

l $\lambda : \text{list}.(0)^*$

n $\eta : \text{word}_4$

↓

assume(l != 0)

l $\lambda : \text{list}.(0)^*$

$\lambda \neq 0$

n $\eta : \text{word}_4$


↓

n = *l ✓

l $\lambda : \text{list}.(0)^*$

$\lambda \neq 0$

n $\eta : \text{list}.(0)^*$



Theorem (Soundness of the abstract semantics)

For all abstract states $s^\# \in \mathbb{S}^\#$ and programs p :

$$(\llbracket p \rrbracket \circ \text{erase_types} \circ \gamma)(s^\#) \subseteq (\text{erase_types} \circ \gamma \circ \llbracket p \rrbracket_S^\#)(s^\#)$$

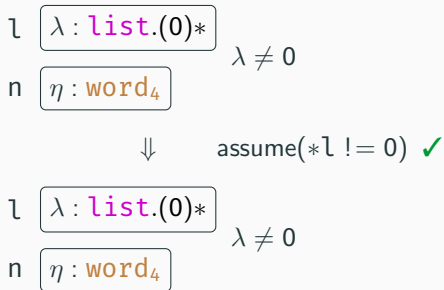
A storeless abstraction is often not enough

\mathfrak{l} $\lambda : \text{list}.(0)^*$ $\lambda \neq 0$
 \mathfrak{n} $\eta : \text{word}_4$

```
type data = word4
type list =
  list.(0)* × data
```

```
if (* $\mathfrak{l}$  != 0) {
   $\mathfrak{n}$  = ** $\mathfrak{l}$ ;
}
```


A storeless abstraction is often not enough



```
type data = word4
type list =
  list.(0)* × data
```

```
if (*l != 0) {
  n = **l;
}
```

A storeless abstraction is often not enough

l $\lambda : \text{list}.(0)*$ $\lambda \neq 0$

n $\eta : \text{word}_4$

\Downarrow $\text{assume}(*\mathsf{l} \neq 0)$ ✓

type $\text{data} = \text{word}_4$
type $\text{list} =$
 $\text{list}.(0)* \times \text{data}$

l $\lambda : \text{list}.(0)*$ $\lambda \neq 0$

n $\eta : \text{word}_4$

\Downarrow $\mathsf{n} = **\mathsf{l}$ ✗

```
if (*l != 0) {  
    n = **l;  
}
```

A storeless abstraction is often not enough

Add points-to
predicates to retain
information about
the heap

λ $\boxed{\lambda : \text{list}.(0)*}$ $\lambda \neq 0$
 n $\boxed{\eta : \text{word}_4}$

```
type data = word4
type list =
  list.(0)* × data
```

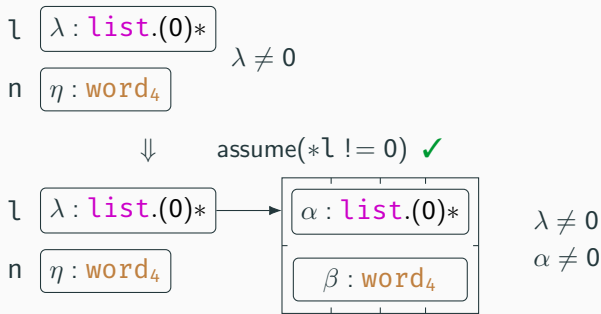
```
if (* $\lambda \neq 0$ ) {
  n = ** $\lambda$ ;
}
```

A storeless abstraction is often not enough

**Add points-to-
predicates to retain
information about
the heap**

```
type data = word4  
type list =  
  list.(0)* × data
```

```
if (*l != 0) {  
  n = **l;  
}
```

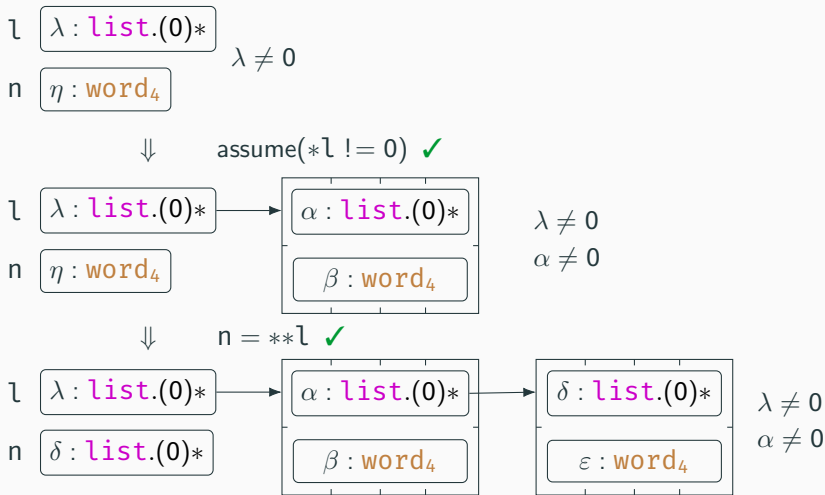


A storeless abstraction is often not enough

**Add points-to-
predicates to retain
information about
the heap**

type **data** = **word**₄
type **list** =
 list.(0)* × **data**

```
if (*l != 0) {  
    n = **l;  
}
```



Temporary violation of structural invariants

\mathfrak{l} $\lambda_0 : \text{word}_4$

n $\eta : \text{list}.(0)^*$

```
type data = word4  
type list =  
  list.(0)* × data
```

```
 $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$   
 $*\mathfrak{l} = n;$   
 $*(\mathfrak{l} + 4) = 0;$ 
```

Temporary violation of structural invariants

\mathfrak{l} $\lambda_0 : \text{word}_4$

n $\eta : \text{list}.(0)^*$

\Downarrow $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$

```
type data = word4
type list =
  list.(0)* × data
```

```
 $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$ 
 $*\mathfrak{l} = n;$ 
 $*(\mathfrak{l} + 4) = 0;$ 
```

Temporary violation of structural invariants

\mathfrak{l} $\lambda_0 : \text{word}_4$

n $\eta : \text{list}.(0)^*$

\Downarrow $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$

\mathfrak{l} $\lambda : \text{list}.(0)^*$

n $\eta : \text{list}.(0)^*$

X

Incorrect: \mathfrak{l} should point to uninitialized memory

```
type data = word4
type list =
  list.(0)* × data
```

```
 $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$ 
 $*\mathfrak{l} = n;$ 
 $*(\mathfrak{l} + 4) = 0;$ 
```


Temporary violation of structural invariants

\mathfrak{l} $\lambda_0 : \text{word}_4$

n $\eta : \text{list}.(0)^*$

\Downarrow $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$

\mathfrak{l} $\lambda : \text{word}_8.(0)^*$

n $\eta : \text{list}.(0)^*$

Correct but very imprecise

```
type data = word4  
type list =  
  list.(0)* × data
```

```
 $\mathfrak{l} = \text{malloc}_{\text{list}}(8);$   
 $*\mathfrak{l} = n;$   
 $*(\mathfrak{l} + 4) = 0;$ 
```

Temporary violation of structural invariants

Represent memory
writes with “staged”
points-to predicates,
delaying type
verification

```
type data = word4  
type list =  
  list.(0)* × data
```

```
l = malloclist(8);  
*l = n;  
*(l + 4) = 0;
```

l $\lambda_0 : \text{word}_4$


n $\eta : \text{list}.(0)^*$

⇓

l = malloc_{list}(8);

l $\lambda : \text{list}.(0)^*$

n $\eta : \text{list}.(0)^*$



Temporary violation of structural invariants

Represent memory
writes with “staged”
points-to predicates,
delaying type
verification

type `data` = `word4`
type `list` =
 `list.(0)* × data`

`l = malloclist(8);`
`*l = n;`
`*(l + 4) = 0;`

`l` `λ0 : word4`


`n` `η : list.(0)*`

⇓

`l = malloclist(8);`

`l` `λ : list.(0)*`

`n` `η : list.(0)*`




⇓

`*l = n; *(l + 4) = 0;`

`l` `λ : list.(0)*`

`n` `η : list.(0)*`



$\beta = 0$

Temporary violation of structural invariants

Represent memory
writes with “staged”
points-to predicates,
delaying type
verification

type `data` = `word4`
type `list` =
 `list.(0)* × data`

```
l = malloclist(8);  
*l = n;  
*(l + 4) = 0;
```

l $\lambda_0 : \text{word}_4$


n $\eta : \text{list}.(0)*$



l = **malloc**_{list}(8);

l $\lambda : \text{list}.(0)*$


n $\eta : \text{list}.(0)*$



*l = n; *(l + 4) = 0;

l $\lambda : \text{list}.(0)*$

n $\eta : \text{list}.(0)*$



$\beta = 0$

Temporary violation of structural invariants

Represent memory
writes with “staged”
points-to predicates,
delaying type
verification

type `data` = `word4`
type `list` =
 `list.(0)* × data`

```
l = malloclist(8);  
*l = n;  
*(l + 4) = 0;
```

l $\lambda_0 : \text{word}_4$


n $\eta : \text{list}.(0)*$



`l = malloclist(8);`

l $\lambda : \text{list}.(0)*$

n $\eta : \text{list}.(0)*$



`*l = n; *(l + 4) = 0;`

l $\lambda : \text{list}.(0)*$

n $\eta : \text{list}.(0)*$



$\beta = 0$

Full analysis: 3,300 lines of OCaml in the Codex
abstract interpretation library

FRAMA-C/CODEX, a module of Frama-C



Software Analyzers

Numerical abstract domain


- Intervals with congruence information
- Bitwise abstraction
- Symbolic equalities and inequalities

Building a machine code analyzer

BINSEC/CODEX, a module of BINSEC




Memory abstraction



"Array of bytes" abstract domain
For the stack and global data

×




Type-based abstract domain
For the rest of memory


Machine code analysis: Incremental inference of control flow

```
001b00: 3d 00 00 00 00
001b04: 00 74 00 48
001b08: eb f8 89 c3
001b0c: f4 00 00 00
```

(a) Binary program



(b) Partial CFG



(c) Invariant CFG

First experimental evaluation: Shape benchmarks

We analyzed all shape benchmarks from [1] and [2], featuring:

- Complex data structure manipulation
- Unstructured sharing (e.g. union-find, graphs)

34 programs, between 9 and 329 lines of code (mean 61)

Target property

Spatial memory safety + Preservation of structural invariants

We evaluate:

- Analysis time
- **Precision** (number of alarms)
- Ease of setup and **automation** (number of annotations)

[1] Li et al. “Shape Analysis for Unstructured Sharing”. In: *SAS*. 2015

[2] Li et al. “Semantic-Directed Clumping of Disjunctive Abstract States”. In: *POPL*. 2017

Analysis process

Generate physical types automatically from C types



Refine types if necessary



Run the analysis



X **✓**

“annotations”

- type predicates
- array lengths



First experimental evaluation: Shape benchmarks

| Benchmark | Annotations gen/ed/pre | | LOC | C | | | O0 | | | O1 | | | O2 | | | O3 | | | |
|--------------------|---------------------------|---|-----|------|------|-------|----------|----------|----------|----------|----------|----------|----|------|----|----|------|----|----|
| | | | | Time | / | ↪ / ↯ | Time/↪/↯ | Time/↪/↯ | Time/↪/↯ | Time/↪/↯ | Time/↪/↯ | Time/↪/↯ | | | | | | | |
| sll-delmin | 11 | 0 | 1 | 25 | 0.27 | 0 | 0 | 0.13 | 0 | 0 | 0.15 | 0 | 0 | 0.15 | 0 | 0 | 0.13 | 0 | 0 |
| sll-delminmax | | | 1 | 49 | 0.30 | 0 | 0 | 0.19 | 0 | 0 | 0.17 | 0 | 0 | 0.17 | 0 | 0 | 0.16 | 0 | 0 |
| psll-bsort | 10 | 0 | 0 | 25 | 0.30 | 0 | 22 | 0.41 | 0 | 3 | 0.25 | 0 | 3 | 0.26 | 0 | 3 | 0.29 | 0 | 3 |
| psll-reverse | | | 0 | 11 | 0.28 | 0 | 2 | 0.10 | 0 | 1 | 0.13 | 0 | 1 | 0.10 | 0 | 1 | 0.10 | 0 | 1 |
| psll-isort | | | 0 | 20 | 0.29 | 0 | 2 | 0.34 | 0 | 1 | 0.34 | 0 | 1 | 0.32 | 0 | 1 | 0.33 | 0 | 1 |
| bstree-find | 12 | 0 | 1 | 26 | 0.27 | 0 | 0 | 0.14 | 0 | 0 | 0.13 | 0 | 0 | 0.15 | 0 | 0 | 0.16 | 0 | 0 |
| gdll-findmin | 25 | 5 | 1 | 49 | 0.50 | 0 | 0 | 0.41 | 0 | 0 | 0.39 | 0 | 0 | 0.41 | 0 | 0 | 0.42 | 0 | 0 |
| gdll-findmax | | | 1 | 58 | 0.55 | 0 | 0 | 0.33 | 0 | 0 | 0.22 | 0 | 0 | 0.21 | 0 | 0 | 0.20 | 0 | 0 |
| gdll-find | | | 1 | 78 | 0.56 | 0 | 0 | 0.15 | 0 | 0 | 0.15 | 0 | 0 | 0.14 | 0 | 0 | 0.14 | 0 | 0 |
| gdll-index | | | 1 | 55 | 0.53 | 0 | 0 | 0.32 | 0 | 0 | 0.33 | 0 | 0 | 0.30 | 0 | 0 | 0.29 | 0 | 0 |
| gdll-delete | | | 1 | 107 | 0.57 | 0 | 2 | 0.16 | 0 | 0 | 0.14 | 0 | 0 | 0.13 | 0 | 0 | 0.13 | 0 | 0 |
| ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ |
| graph-nodelisttrav | 23 | 0 | 1 | 12 | 0.20 | 0 | 0 | 0.10 | 0 | 0 | 0.10 | 0 | 0 | 0.10 | 0 | 0 | 0.11 | 0 | 0 |
| graph-path | | | 1 | 19 | 0.21 | 0 | 14 | 0.15 | 0 | 5 | 0.16 | 0 | 0 | 0.14 | 0 | 0 | 0.16 | 0 | 0 |
| graph-pathrand | | | 1 | 25 | 0.22 | 0 | 10 | 0.13 | 0 | 0 | 0.21 | 0 | 0 | 0.12 | 0 | 0 | 0.11 | 0 | 0 |
| graph-edgeadd | | | 1 | 15 | 0.27 | 0 | 2 | 0.12 | 0 | 1 | 0.11 | 0 | 1 | 0.10 | 0 | 1 | 0.10 | 0 | 1 |
| graph-nodeadd | | | 1 | 15 | 0.26 | 0 | 2 | 0.10 | 0 | 1 | 0.08 | 0 | 1 | 0.09 | 0 | 1 | 0.10 | 0 | 1 |
| graph-edgededelete | | | 1 | 11 | 0.20 | 0 | 2 | 0.10 | 0 | 1 | 0.10 | 0 | 0 | 0.10 | 0 | 0 | 0.11 | 0 | 0 |
| graph-edgeiter | | | 1 | 22 | 0.23 | 0 | 0 | 0.13 | 0 | 0 | 0.11 | 0 | 0 | 0.12 | 0 | 0 | 0.12 | 0 | 0 |
| uf-find | 33 | 3 | 1 | 11 | 0.31 | 0 | 24 | 0.07 | 0 | 6 | 0.09 | 0 | 0 | 0.08 | 0 | 0 | 0.07 | 0 | 0 |
| uf-merge | | | 1 | 17 | 0.34 | 0 | 50 | 0.13 | 0 | 7 | 0.18 | 0 | 0 | 0.18 | 0 | 0 | 0.15 | 0 | 0 |
| uf-make | | | 0 | 9 | 0.31 | 0 | 4 | 0.05 | 0 | 3 | 0.06 | 0 | 3 | 0.07 | 0 | 3 | 0.06 | 0 | 3 |
| Total verified | | | | | | 30 | 13 | | 30 | 16 | | 30 | 21 | | 30 | 21 | | 30 | 21 |

The annotation effort is low

Annotation effort varies between **0** and **12 lines** (on average **3.2 lines**).

| | min. | mean | median | max. |
|-----------------------|------|-------|--------|-------|
| Annotation/code ratio | 0 % | 3.2 % | 2.7 % | 7.8 % |

Precision

| | C | binary -00 | binary -01 | binary -02 | binary -03 |
|---------------------|----------------|----------------|----------------|----------------|----------------|
| Property verified ✓ | 30 / 34 | 30 / 34 | 30 / 34 | 30 / 34 | 30 / 34 |

Points-to predicates are necessary

Without points-to predicates:

| | C | binary -00 | binary -01 | binary -02 | binary -03 |
|---------------------|---------|------------|------------|------------|------------|
| Property verified ✓ | 13 / 34 | 16 / 34 | 21 / 34 | 21 / 34 | 21 / 34 |

The analysis time is predictable

All analyses complete in **less than 1 s**.

Comparison with shape analyses from [2]

| | Base shape analysis | Guided clumping [2] | This work |
|------------------------|------------------------|------------------------|------------------|
| javl-insert (baseline) | X | 1.84 | 0.43 |
| javl-insert-32× | X | 129.9 | 40 |

Run times (s)

Overview: Kernel verification using BINSEC/CODEX

We propose a method based on BINSEC/CODEX to verify **APE** and **ARTE** on **embedded kernels**.

- **Automated**
 - Abstract interpretation *infers* kernel invariants
 - APE is an implicit property (no specification to write)
- **Comprehensive**
 - *Machine code* verification on the kernel executable
- **Generic**
 - *Parameterized* verification by *abstracting* task-specific data with *types*
- **Practical**
 - Comprehensive evaluation on challenging case studies
unmodified version of ASTERIOS RTK, 96 variants of EducRTOS

Definition (Implicit property)


An implicit property is a property that does not depend on a particular program.

Example

Absence of run-time errors (**ARTE**) is an implicit property.

Verifying an implicit property does not require to **write a specification**.


Abstraction of user code



Alternation of **user code** and **kernel runtime**.

The **user code** is unknown

Abstraction of user code



Alternation of **user code** and **kernel runtime**.

The **user code** is unknown

⇒ We abstract it by “arbitrary sequences of instructions”
(whose execution is permitted by the hardware).

Hardware protection mechanisms

- Memory protection
- Hardware privilege level

Absence of Privilege Escalation is an implicit property

Theorem

*If the system satisfies a non-trivial invariant,
then no privilege escalation is possible on that system.*


Proof.

If the system fails to self-protect, the empowered attacker can reach any state. □

⇒ APE can be verified without writing a specification.

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

○ ————— $cur \in \{0xa7\}$, $mpu \in \{0xae\}$

Example in-context analysis

Initial state:



Task *cur;


```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

○ ——— cur ∈ {0xa7}, mpu ∈ {0xae}

○ ——— cur ∈ {0xa7}, mpu ∈ {0xae}

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

○ ————— $cur \in \{0xa7\}$, $mpu \in \{0xae\}$
○ ————— $cur \in \{0xa7\}$, $mpu \in \{0xae\}$
○ ————— $cur \in \{0xa2\}$, $mpu \in \{0xae\}$

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```


Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

○ ————— $cur \in \{0xa7\}, mpu \in \{0xae\}$
○ ————— $cur \in \{0xa7\}, mpu \in \{0xae\}$
○ ————— $cur \in \{0xa2\}, mpu \in \{0xae\}$
○ ————— $cur \in \{0xa2\}, mpu \in \{0xae\}$
○ ————— $cur \in \{0xa2\}, mpu \in \{0xae\}$

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

Example in-context analysis

Initial state:




Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

Example in-context analysis

Initial state:



Task *cur;

```
void handle_timer() {  
    /* Save task context */  
    cur->ctx = interrupted_ctx;  
    /* Schedule next task */  
    cur = cur->next;  
    /* Load new memory protection */  
    mpu = cur->mem_table;  
    /* Give control to new task */  
    switch_context(&cur->ctx);  
}
```

user code#

BINSEC/CODEX can verify APE and ARTE of such small kernels with 0 lines of annotations.

Shortcomings of in-context analyses

In-context analysis is:

- *Not generic*: Cannot analyze kernel independently from the applications
- *Not scalable*: 1000 tasks \implies 1000 addresses to enumerate.

Key idea

We use the type-based abstract domain to abstract task data.

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

$\alpha \neq 0$ $\beta \neq 0$

```
void handle_timer() {  ○———— cur [α: Task.(0)*]  mpu [β: Memory_table.(0)*]
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

```
Task *cur;
```

$\alpha \neq 0$ $\beta \neq 0$

```
void handle_timer() {  ○———— cur [α: Task.(0)*]  mpu [β: Memory_table.(0)*]
  /* Save task context */
  cur→ctx = interrupted_ctx;  ○— cur [α: Task.(0)*]  mpu [β: Memory_table.(0)*]
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\beta \neq 0$ $\delta \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

| | | | |
|-----|------------------------------|-----|--------------------------------------|
| cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| cur | $\delta : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\beta \neq 0$ $\delta \neq 0$ $\varepsilon \neq 0$

```
void handle_timer() {  ○————— cur  $\alpha : \text{Task}.(0)^*$     mpu  $\beta : \text{Memory\_table}.(0)^*$ 
  /* Save task context */
  cur→ctx = interrupted_ctx;  ○— cur  $\alpha : \text{Task}.(0)^*$     mpu  $\beta : \text{Memory\_table}.(0)^*$ 
  /* Schedule next task */
  cur = cur→next;  ○————— cur  $\delta : \text{Task}.(0)^*$     mpu  $\beta : \text{Memory\_table}.(0)^*$ 
  /* Load new memory protection */
  mpu = cur→mem_table;  ○————— cur  $\delta : \text{Task}.(0)^*$     mpu  $\varepsilon : \text{Memory\_table}.(0)^*$ 
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\beta \neq 0$ $\delta \neq 0$ $\varepsilon \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

| | | | | |
|--------|-----|------------------------------|-----|--|
| ○————— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○—— | cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| ○————— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |
| ○—— | cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |

Example parameterized analysis

```
type Memory_table = {x: int | x > kernel_last_addr} × int
type Context = ... × {privilege: int | privilege = 0} × ...
type Task = {x: Memory_table.(0)* | x ≠ 0} × Context × {x: Task.(0)* | x ≠ 0}
```

Task *cur;

$\alpha \neq 0$ $\beta \neq 0$ $\delta \neq 0$ $\varepsilon \neq 0$

```
void handle_timer() {
  /* Save task context */
  cur→ctx = interrupted_ctx;
  /* Schedule next task */
  cur = cur→next;
  /* Load new memory protection */
  mpu = cur→mem_table;
  /* Give control to new task */
  switch_context(&cur→ctx);
}
```

| | | | |
|-----|------------------------------|-----|--|
| cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| cur | $\alpha : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| cur | $\delta : \text{Task}.(0)^*$ | mpu | $\beta : \text{Memory_table}.(0)^*$ |
| cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |
| cur | $\delta : \text{Task}.(0)^*$ | mpu | $\varepsilon : \text{Memory_table}.(0)^*$ |


user code[#]

Differentiated handling of boot and runtime code

- Type-based analysis verifies the **preservation** of the invariant
- But the boot code **establishes** that invariant

Based on this, we

1. Perform a *parameterized* analysis of the *runtime*
2. And an *in-context* analysis of the boot code
3. Check that the state after boot matches the invariant.



Experimental evaluation: Real-life effectiveness

Case study 1: ASTERIOS

- Industrial microkernel used in industrial settings
- Version: port to an ARM quad-core
- 329 functions, ~10,000 instructions
- Protection using *page tables*.

2 versions

- **BETA** version: 1 vulnerability
- **v1** version: vulnerability fixed

Specific = restriction on stack sizes

| | | Generic annotations | | Specific annotations | |
|-----------------------|-----------|---------------------------------------|---------------|--------------------------------------|-----|
| # shape annotations | generated | 1057 | | | |
| | manual | 57 (5.11%) | | 58 (5.20%) | |
| Kernel version | | BETA | v1 | BETA | v1 |
| invariant computation | status | ✓ | ✓ | ✓ | ✓ |
| | time (s) | 647 | 417 | 599 | 406 |
| # alarms in runtime | | 1 true error 2 false alarms | 1 false alarm | 1 true error 1 false alarm | 0 ✓ |
| user tasks checking | status | ✓ | ✓ | ✓ | ✓ |
| | time (s) | 32 | 29 | 31 | 30 |
| Proves APE? | | N/A | ~ | N/A | ✓ |

Proved APE and ARTE in 430 s.
58 lines of annotations.

Case study 2: EducRTOS

- Small academic OS developed for teaching purposes
- Both separation kernel and real-time OS, dynamic thread creation
- 1,200 x86 instructions.
- Protection by *segmentation*.

Proved APE and ARTE on 96 variants.

Varying parameters:

- compiler (GCC/Clang), optimization flags
- scheduling algorithm (EDF/FP) dynamic thread creation (on/off)
- ...


Verification time: from 1.6 s to 73 s.

14 lines of annotations.

Experimental evaluation: Automation and Scalability

We compare

- *fully automated in-context analysis vs parameterized analysis* (12 lines of annotations)
- for a simple variant of EducRTOS
- with varying numbers of tasks.



Time and space complexity of *parameterized* analysis is *almost linear*
In-context verification is *quadratic*

- A **type-based** memory abstract domain
 - Points-to predicates improve precision without disjunctions
- Two **low-level program analyses** in Binsec and Frama-C
 - Verify **structural invariants** and **spatial memory safety**
 - Less precise but **more easily applicable** than shape analyses
- A method to verify **absence of run-time errors** and **absence of privilege escalation** on embedded kernels.

- **Collaboration** with other proof methods
 - Other methods verify more properties but establishing invariants is costly
 - Use facts inferred by our analysis as input to other verification tools
- Enrich the type system to verify **stronger properties**
 - Full dependent types (dynamic-size arrays, non-interference)
 - Combination with **separation logic** to specify local separation constraints